

# Resources and tasks management in ES

(up to/from here: session 19/20)

Amir Mahdi Hosseini Monazzah

Room 332,  
School of Computer Engineering,  
Iran University of Science and Technology,  
Tehran, Iran.

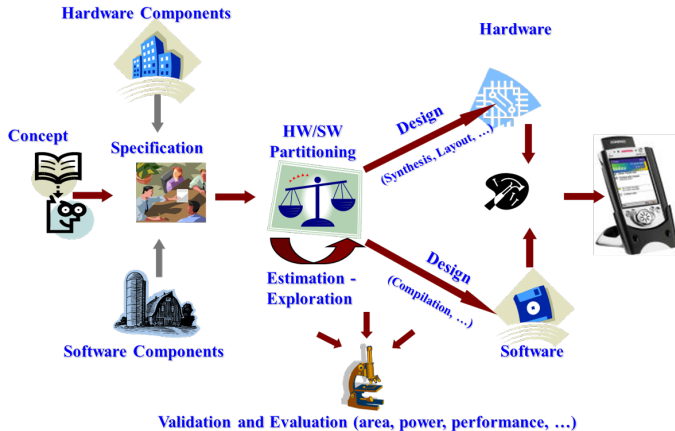
*monazzah@iust.ac.ir*

Fall 2020

# Outline

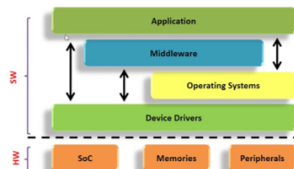
- Preliminaries on scheduling
  - A deeper view on ES software
  - Definitions
  - Classifications of scheduling algorithms
- Real-time scheduling
  - Definitions
  - Different types of real-time scheduling algorithms
  - Uniprocessor real-time scheduling
    - Equal arrival times algorithms
    - Different arrival times
  - Multi-processor real-time scheduling

# Embedded system design flow



# Reuse of standard software components

- Knowledge from previous designs to be made available in the form of intellectual property (IP, for SW and HW).
  - Operating systems (real-time)
  - Middleware
  - Real-time data bases
  - Standard software (MPEG-x, GSM-kernel, ...)
- Includes standard approaches for scheduling (requires knowledge about execution times)



# Time-triggered systems

- Entirely time-triggered system
  - The temporal control structure of all tasks is established a priori by off-line support-tools
  - Temporal control structure is encoded in a Task-Descriptor List (TDL) that contains the cyclic schedule for all activities of the node
  - This schedule considers the required precedence and mutual exclusion relationships among the tasks such that an **explicit coordination** of the tasks by the operating system at run time is **not necessary**

# Time-triggered systems

- The dispatcher is activated by the synchronized clock tick
  - It looks at the TDL, and then performs the action that has been planned for this instant [Kopetz]
  - The only practical means of providing predictability
- It can be easily checked if timing constraints are met
  - Response to sporadic events may be poor

# Worst case execution time

- **Definition: The worst case execution time (WCET) is an upper bound on the execution times of tasks**
  - Computing such a bound is undecidable
    - Possible for programs without recursion and finite loops
      - **Pipeline hazards, interrupts, caches** → **serious overestimates**
- **Approaches:**
  - For hardware: Typically requires hardware synthesis
  - For software: Requires availability of machine programs; complex analysis

# Average execution time

- Estimated cost and performance values
  - Difficult to generate sufficiently precise estimates; Balance between run-time and precision
- Accurate cost and performance values
  - Can be done with normal tools (such as compilers)
    - As precise as the input data is

Which timing metric is more important in real-time ES?



# Schedulability

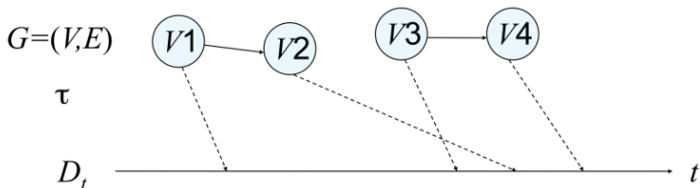
- A set of tasks is said to be schedulable under a given set of constraints, if a schedule exists for that set of tasks and constraints
  - Exact tests are NP-hard in many situations
  - Sufficient tests: Sufficient conditions for guaranteeing a schedule are checked
  - Necessary tests: Checking necessary conditions. Can be used to show that no schedule exists
    - Always not possible to prove even if no schedule exists

# Scheduling classifications

- Centralized and distributed scheduling
  - Multiprocessor scheduling either locally on one or distributed on several processors
- Mono- and multi-processor scheduling
  - Simple scheduling algorithms handle single processors, more complex algorithms handle multiple processors
- Online (dynamic) - and offline (static) scheduling:
  - Online scheduling is done at run-time based on the information about the tasks arrived so far
  - Offline scheduling assumes prior knowledge about arrival times, execution times, and deadlines

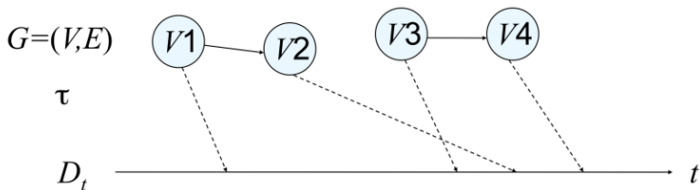
# What is real-time scheduling?

- Assume that we have a task graph  $G=(V,E)$
- A schedule of  $G$  is a mapping,  $V \rightarrow T$ , of a set of tasks  $V$  to start times from domain  $T$



# What is real-time scheduling?

- Schedules have to respect a set of constraints, such as resource, dependency, and **deadlines**
- Scheduling is the process of finding such a mapping
- During the design of embedded systems, scheduling has to be performed several times
  - Early rough scheduling as well as late precise scheduling



# Simple tasks

- Tasks without any inter-process communication are called simple tasks (S-tasks)
  - S-tasks can be either **ready** or **running**
- API of an S-task in a TT system: Two OS Calls
  - TERMINATE TASK and ERROR
    - The TERMINATE TASK system call is executed whenever the task has reached its termination point
    - In case of an error that cannot be handled within the application task, the task terminates its operation with the ERROR system call [Kopetz, 1997]

# Cost functions

- Cost function: Different algorithms aim at minimizing different functions
  - For example **Maximum lateness** is one of the main metrics that are used in cost functions
- **Definition:** Maximum lateness is defined as the difference between the completion time and the deadline, maximized over all tasks
  - Maximum lateness is **negative** if all tasks complete before their deadline

# Basic parameters in real-time scheduling (from here: session 21)

- Arrival time  $a_i$ 
  - Time when a task becomes ready for execution
- Computation time  $c_i$ 
  - Time necessary for completion of a task
- Absolute deadline  $d_i$ 
  - Time before which a task should be finished
- Finishing time  $f_i$ 
  - Time at which a task finishes its execution

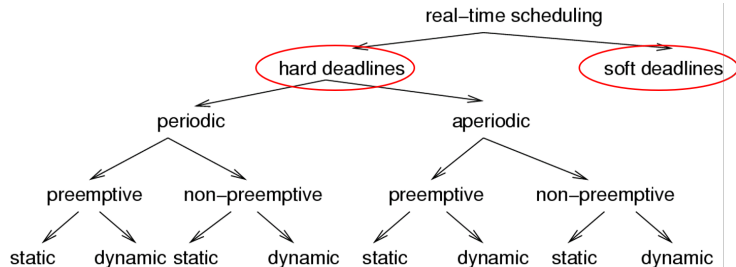
# Basic parameters in real-time scheduling (up to here: session 20)

- Derived Parameters

- Relative deadline  $D_i = d_i - a_i$
- Response time  $R_i = f_i - a_i$
- Lateness  $L_i = f_i - d_i$  delay of a task (can be negative)
- Slack time (laxity)  $s_i = d_i - a_i - c_i$ 
  - Maximum time a task can be delayed on its activation to complete within deadline

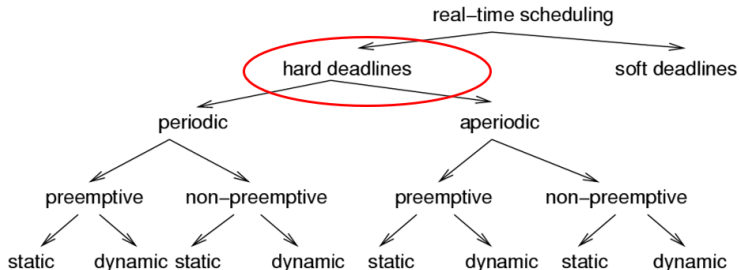


# Classification of real-time scheduling algorithms

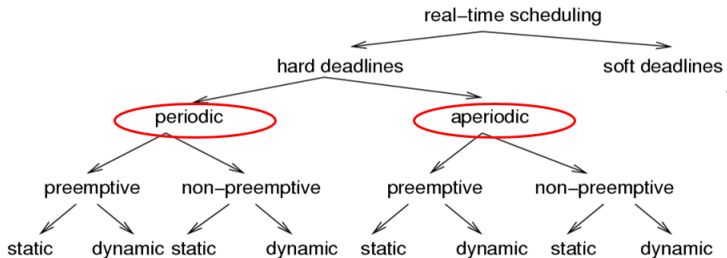


# Hard and soft deadlines

- **Definition:** A time-constraint (deadline) is called hard if not meeting that constraint could result in a catastrophe [Kopetz, 1997]
  - All other time constraints are called soft
- We will focus on hard deadlines!

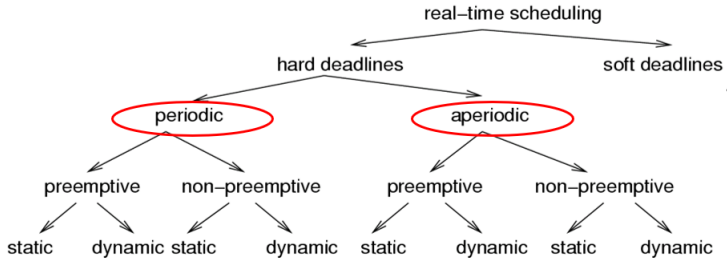


# Periodic and aperiodic tasks



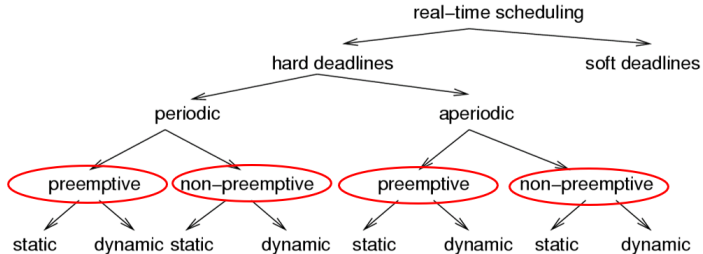
- **Definition:** Tasks which must be executed once every  $p$  units of time are called **periodic tasks**
  - $p$  is called their period
  - Each execution of a periodic task is called a job
  - All other tasks are called aperiodic

# Periodic and aperiodic tasks



- **Definition:** Aperiodic tasks requesting the processor at unpredictable times are called sporadic if there is a minimum separation between the times at which they request the processor

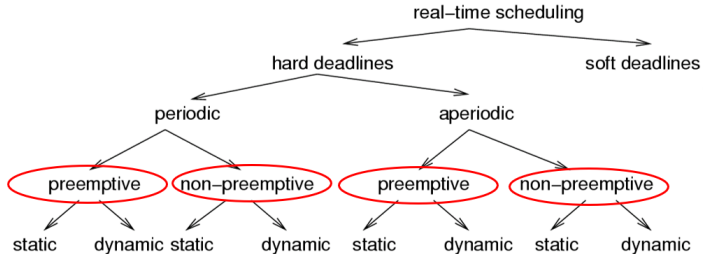
# Preemptive and non-preemptive scheduling



- **Definition: Preemptive and non-preemptive scheduling**

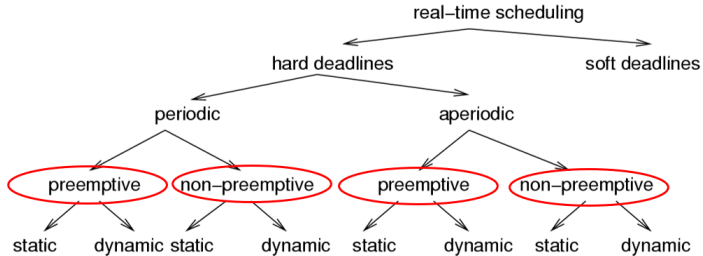
- Non-preemptive schedulers are based on the assumption that tasks are executed until they are done

# Preemptive and non-preemptive scheduling



- For non-preemptive scheduler
  - The response time for external events may be quite long if some tasks have a large execution time

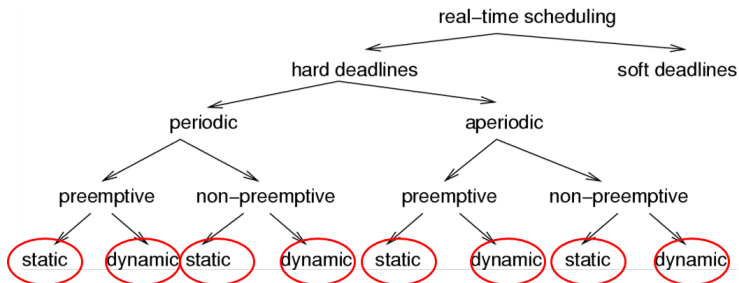
# Preemptive and non-preemptive scheduling



- Preemptive schedulers have to be used if some tasks have long execution times or if the response time for external events is required to be short

# Static and dynamic scheduling

- **Definition: Dynamic scheduling**
  - **Processor allocation decisions (scheduling) done at run-time**

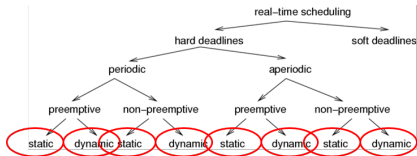





# Static and dynamic scheduling

## • Definition: Static scheduling

- Processor allocation decisions (scheduling) done at design-time
  - Dispatcher allocates processor when interrupted by a timer
  - The timer is controlled by a table generated at design time

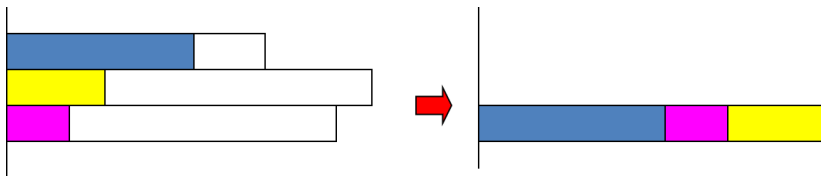


Time	Action	WCET
10	start T1	12
17	send M5	
22	stop T1	
38	start T2	20
47	send M3	

→ 

# Earliest Due Date (EDD)

- Given a set of  $n$  independent tasks, any algorithm that executes the tasks in order of non-decreasing deadlines is optimal with respect to minimizing the maximum lateness
  - Proof: See Dertouzos, M. L.: "Control Robotics: the Procedural Control of Physical Processes", Information Processing 74, North-Holland Publishing Company, 1974
- EDD requires all tasks to be sorted by their deadlines



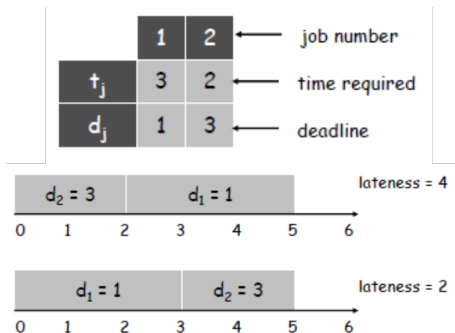
# In-class assignment: Question

- Consider the following job's information table
  - Draw all possible scheduling Gantt chart
  - Which one is better in terms of minimizing maximum lateness?

	1	2	← job number
$t_j$	3	2	← time required
$d_j$	1	3	← deadline

# In-class assignment: Answer

- Consider the following job's information table
  - Draw all possible scheduling Gantt chart
  - Which one is better in terms of minimizing maximum lateness?



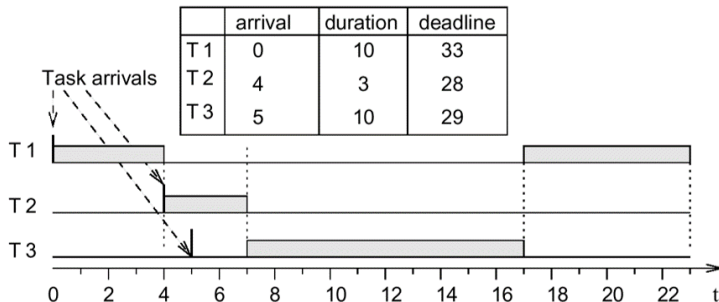
# Earliest Deadline First (EDF)

- Theorem [Horn74]: Given a set of  $n$  independent tasks with arbitrary arrival times, any algorithm that at any instant executes the task with the **earliest absolute deadline** among all the ready tasks is optimal with respect to minimizing the maximum lateness

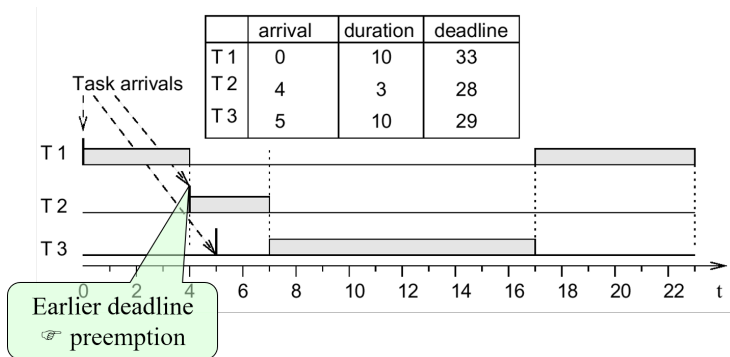
# Earliest Deadline First (EDF)

- Each time a new ready task arrives, it is inserted into a queue of ready tasks, sorted by their deadlines
  - If a newly arrived task is inserted at the head of the queue, the currently executing task is preempted
    - Preemption potentially reduces lateness

# Earliest Deadline First (EDF)

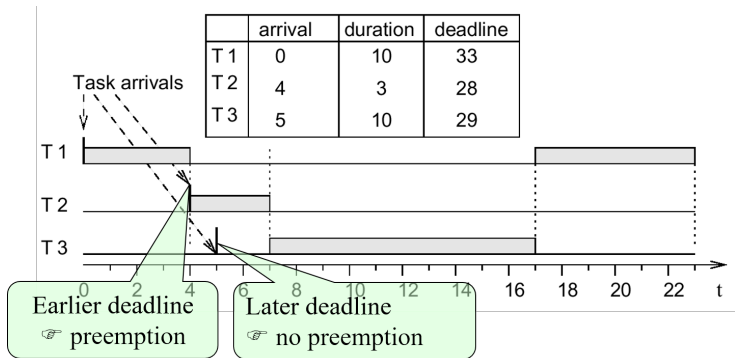


# Earliest Deadline First (EDF)





# Earliest Deadline First (EDF)

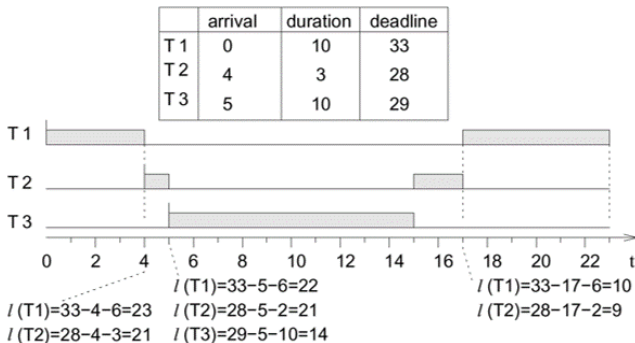


# Scheduling with no precedence constraints

- Let  $T_i$  be a set of tasks and
  - $c_i$  be the execution time of  $T_i$
  - $d_i$  be the **deadline interval**, that is, the time between  $T_i$  becoming available and the time until which  $T_i$  has to finish execution
  - $l_i$  be the **laxity** or **slack**, defined as  $l_i = d_i - c_i$

# Least Laxity First (LLF) or Least Slack Time First (LSF)

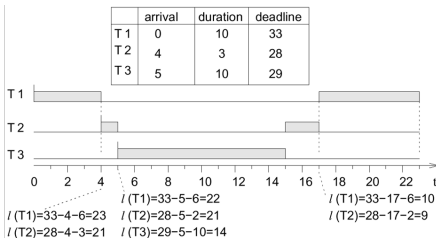
- Priorities = Decreasing function of the laxity
  - The less laxity, the higher the priority
  - Dynamically changing priority
  - Preemptive



# Least Laxity First (LLF) or Least Slack Time First (LSF)

- Priorities = Decreasing function of the laxity
  - The less laxity, the higher the priority
  - Dynamically changing priority
  - Preemptive

- ✓ Requires calling the scheduler periodically, and to recompute the laxity. Overhead for many calls of the scheduler and many context switches.
- ✓ Detects missed deadlines early



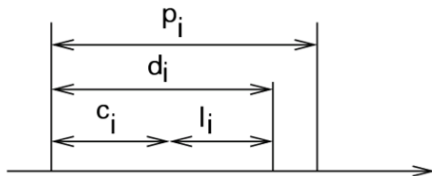
# LLF/LSF properties

- LLF/LSF is also an optimal scheduling for uni-processor systems
  - BUT... uses dynamic priorities
    - Therefore cannot be used with a fixed priority OS
      - Fixed-priority preemptive scheduling is a scheduling system commonly used in real-time systems
- LLF/LSF scheduling requires the knowledge of the execution time
  - May not know this in advance!

# Periodic scheduling (from here: session 22)

- Let

- $p_i$  be the period of task  $T_i$ ,
- $c_i$  be the execution time of  $T_i$ ,
- $d_i$  be the deadline interval, that is, the time between a job of  $T_i$  becoming available and the time after which the same job  $T_i$  has to finish execution
- $l_i$  be the **laxity** or **slack**, defined as  $l_i = d_i - c_i$



# Periodic scheduling (up to here: session 21)

- Accumulated utilization
  - Accumulated execution time divided by period

$$\mu = \sum_{i=1}^n \frac{c_i}{p_i}$$

Necessary condition for schedulability  
(with  $m$ =number of processors):

$$\mu \leq m$$

# Rate Monotonic (RM)

- Well-known technique for scheduling independent periodic tasks [Liu, 1973]
- Assumptions:
  - All tasks that have hard deadlines are periodic
  - All tasks are independent
  - $d_i = p_i$ , for all tasks
  - $c_i$  is constant and is known for all tasks
  - The time required for context switching is negligible



# Rate Monotonic (RM)

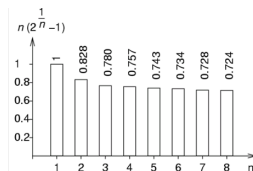
- RM schedulability condition
  - For a single processor with  $n$  tasks, the following equation must hold for the accumulated utilization  $\mu$

$$\mu = \sum_{i=1}^n \frac{c_i}{p_i} \leq n(2^{1/n} - 1)$$

# Rate Monotonic (RM)

- The priority of a task is a monotonically decreasing function of its period
  - Low period: High priority
- At any time, a highest priority task among all those that are ready for execution is allocated
- If all assumptions are met, schedulability is guaranteed

Maximum utilization as a function of the number of tasks  $\Rightarrow$



# Example: RM-generated schedule

- T1 preempts T2 and T3
- T2 and T3 do not preempt each other



# Case of failing RM scheduling

Task 1: period 5, execution time 2

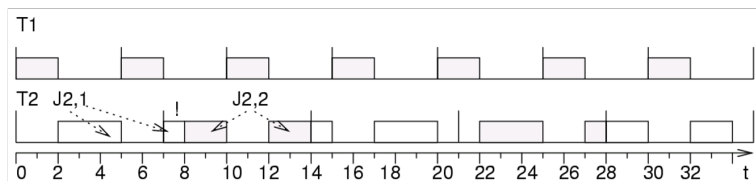
Task 2: period 7, execution time 4

$$\mu = 2/5 + 4/7 = 34/35 \approx 0.97$$

$$2(2^{1/2} - 1) \approx 0.828$$

$$\mu = \sum_{i=1}^n \frac{c_i}{p_i} \leq n(2^{1/n} - 1)$$

**Not enough idle time**



# Case of failing RM scheduling

Task 1: period 5, execution time 2

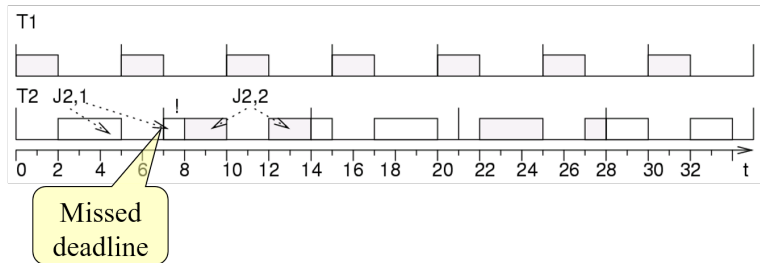
Task 2: period 7, execution time 4

$$\mu = 2/5 + 4/7 = 34/35 \approx 0.97$$

$$2(2^{1/2} - 1) \approx 0.828$$

$$\mu = \sum_{i=1}^n \frac{c_i}{p_i} \leq n(2^{1/n} - 1)$$

**Not enough idle time**



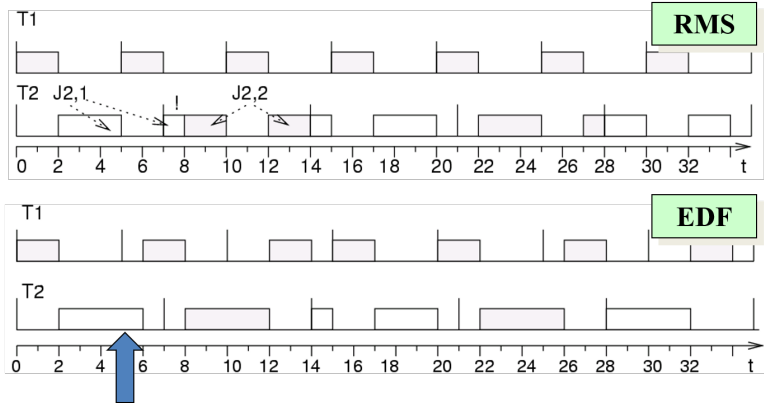
# Properties of RM scheduling

- RM scheduling is based on static priorities
  - This allows RM scheduling to be used in standard OS
    - Such as Windows NT
  - A huge number of variations of RM scheduling exists
- In the context of RM scheduling, many formal proofs exist.
- The **idle capacity** is not required if the period of all tasks is a multiple of the period of the highest priority task
  - Necessary condition for schedulability:  $\mu \leq 1$

# EDF in periodic scheduling

- EDF can also be applied to periodic scheduling
- EDF optimal for every period
  - Optimal for periodic scheduling
    - Trivially!
  - EDF must be able to schedule the example in which RMS failed
- EDF requires dynamic priorities
  - EDF cannot be used with a standard operating system just providing static priorities

# Comparison of EDF and RMS



T2 not preempted, due to its earlier deadline



# Scheduling without preemption

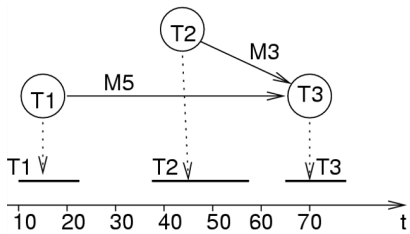
- Optimal schedules may leave processor idle to finish tasks with early deadlines arriving late
  - Knowledge about the future is needed for optimal scheduling algorithms
    - No online algorithm can decide whether or not to keep idle

# Scheduling without preemption

- EDF is optimal among all scheduling algorithms not keeping the processor idle at certain times
- If arrival times are known a priori, the scheduling problem becomes NP-hard in general
  - Branch and bound techniques are typically used

# Scheduling with precedence constraints

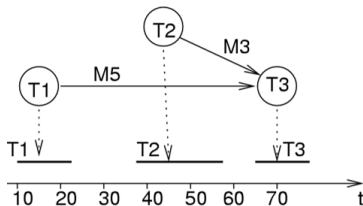
- Task graph and possible schedule



Schedule can be stored in a table  
(can be used by dispatcher/OS)

# Synchronous arrival times

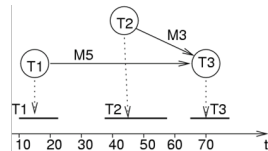
- Optimal algorithm for minimum latency
  - Latest Deadline First (LDF)
- LDF [Lawler, 1973]: Generation of total order compatible with the partial order described by the task graph (LDF performs a **topological sort**)



# Synchronous arrival times

- LDF reads the task graph from tail to head and inserts tasks with no successors into a queue. It then repeats this process, putting tasks whose successor have all been selected into the queue
  - At run-time, the tasks are executed in from head to tail
  - LDF is non-preemptive and is optimal for uni-processors

## Latest Deadline First (LDF)



# Asynchronous arrival times

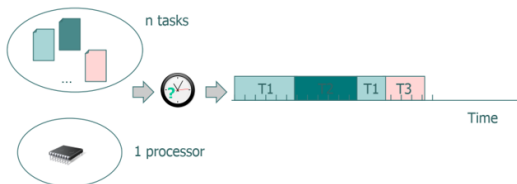
- This case can be handled with a modified EDF algorithm
  - The key idea is to transform the problem from a given set of dependent tasks into a set of independent tasks with different timing parameters [Chetto90]
- This algorithm is optimal for uni-processor systems
- If preemption is not allowed, the heuristic algorithm developed by [Stankovic and Ramamritham 1991] can be used

# Sporadic tasks

- If sporadic tasks were connected to interrupts, the execution time of other tasks would become very unpredictable
  - Introduction of a sporadic task server, periodically checking for ready sporadic tasks
    - Sporadic tasks are essentially turned into periodic tasks

# Introduction (from here: session 23)

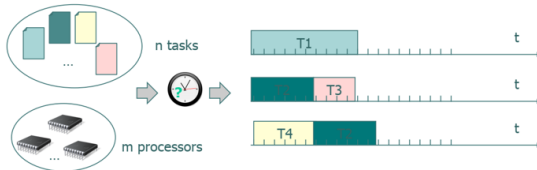
- Mono-processor scheduling: One-dimension problem
  - **Temporal** organization





# Introduction

- Multi-processor (multi-core) scheduling: Two dimension problem
  - **Temporal** organization +
  - **Spatial** organization
    - On which processor execute every task?

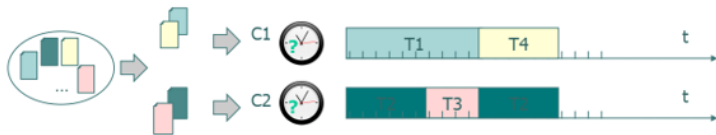


# Classification (up to here: session 22)

- **Partitioned** scheduling
  - Each of the two dimensions is dealt with separately
- **Global** scheduling
  - Temporal and spatial dimensions are deal with jointly
- **Semi-partitioned** scheduling
  - Hybrid

# Classification: partitioned scheduling

- Each of the two dimensions is dealt with separately
  - Spatial organization: the  $n$  tasks are partitioned onto the  $m$  cores
    - No task migration at run-time
  - Temporal organization: Mono-processor scheduling is used on each core



# Classification: partitioned scheduling

- Two points of view
  - Number of processors to be determined: Optimization problem (bin-packing problem)
    - Bin = task, size = utilization (or other expression obtained from the task temporal parameters)
    - Boxes = processors, size = ability to host tasks
  - Fixed number of processors: search problem (knapsack problem)
- Both problems are NP-hard

# Classification: partitioned scheduling

- Optimal mono-processor scheduling strategies: XX
  - RM, DM (Deadline Monotonic: Like rate monotonic but consider **task's deadlines** instead of task's periods)
  - EDF, LLF (see uni-processor scheduling section)
- Bin-packing heuristics: YY
  - FF: First-Fit
  - BF: Best-Fit
  - WF: Worst-Fit
  - NF: Next-Fit
  - FFD, BFD, WFD:  
First/Best/Worst-Fit Decreasing

Partitioning algorithms  
XX-YY

# Classification: partitioned scheduling

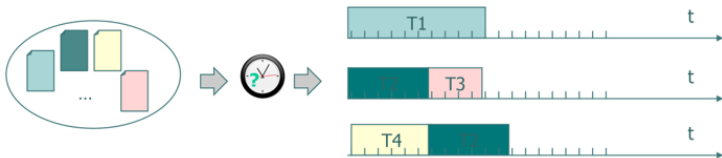
- Benefits
  - Implementation: local schedulers are independent
  - No migration costs
  - Direct reuse of mono-processor schedulability tests
  - Isolation between processors in case of overload
- Limits

# Classification: partitioned scheduling

- Benefits
- Limits
  - Rigid: suited to static configurations
  - NP-hard task partitioning
  - Largest utilization bound for **any** partitioning algorithm [Andersson, 2001]
    - $m+1$  tasks of execution time " $1 + \epsilon$ " and period 2:  $\frac{m+1}{2}$

# Classification: global scheduling

- Temporal and spatial dimensions are dealt with jointly
  - Global unique scheduler and run queue
  - At each scheduling point, the scheduler decides when and where to schedule at most  $m$  tasks
  - Task migration allowed



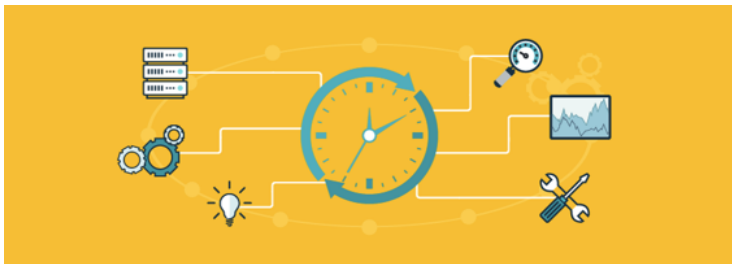


# Classification: global scheduling

- Benefits
  - Suited to dynamic configurations
  - Dominates all other scheduling policies
    - If we consider unconstrained migrations + dynamic priorities
  - Optimal schedulers exist
  - Overloads/underloads spread on all processors
- Drawbacks
  - System overheads: migrations, mutual exclusion for sharing the run queue

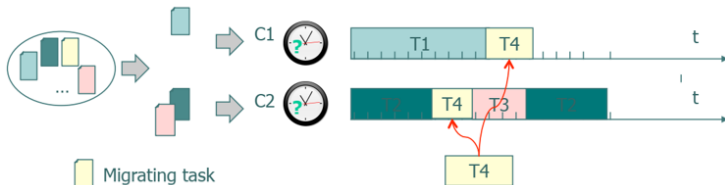
# Classification: global scheduling

- Global RM/DM/EDF (preemptive): definition
  - Task priorities assigned according to RM/DM/EDF
  - Scheduling algorithm: The  $m$  higher priority tasks are executed on the  $m$  processors



# Classification: semi-partitioned scheduling

- Partitioned scheduling as far as possible
- Some statically determined tasks may migrate
  - Constraint: Migrating tasks (T4 on the example) must execute on a single processor at a time



# Terminology

- Priorities
  - Fixed per task (FPT)
  - Fixed per job (FPJ)
  - Dynamic per job (DPJ)

# Overview of global scheduling policies (from here: session 24)

- Assumptions

- Tasks

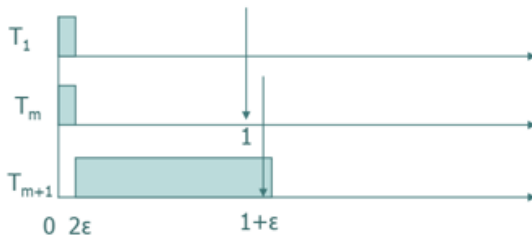
- Periodic tasks ( $p_i$ )
- Implicit deadlines ( $d_i = p_i$ )
- Synchronous tasks
- Independent tasks
- A single job of a task can be active at a time

- Architecture

- Identical processors
- Costs are neglected (preemption, migration, scheduling policy)

# Scheduling anomalies

- Dhall's effect [Dhall Liu, 1978]
  - Periodic task sets with utilization close to 1 are unschedulable using global RM/EDF
  - $n=m+1$  ,  $p_i=1$  ,  $c_i=2\epsilon$  ,  $u_i=2\epsilon$  for all  $1 \leq i \leq m$
  - $P_{m+1}=1+\epsilon$  ,  $C_{m+1}=1$  ,  $u_{m+1}=\frac{1}{1+\epsilon}$
  - Task  $m+1$  misses its deadline although  $U$  very close to 1
  - We assumed  $m$  processor cores

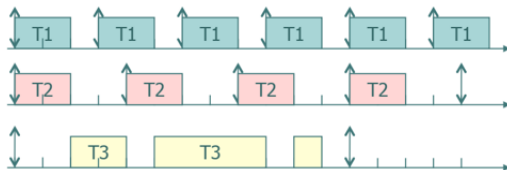


# Scheduling anomalies (up to here: session 23)

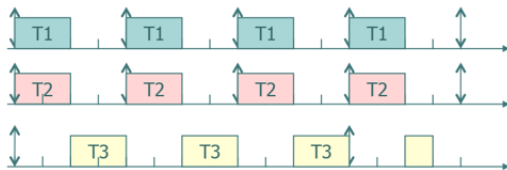
- For G-RM, there may be situations in which schedules exist for a certain task system, but deadlines are violated if periods are extended [Anderson, 2003]
  - $n = 3$  ,  $m = 2$  ,  $(P_1 = 3, C_1 = 2), (P_2 = 4, C_2 = 2), (P_3 = 12, C_3 = 7)$
  - Schedulable under global RM
  - If  $P_1$  is increased to 4 and priorities stay the same, T3 misses its deadline

# Scheduling anomalies

- $(P_1 = 3, C_1 = 2), (P_2 = 4, C_2 = 2), (P_3 = 12, C_3 = 7)$



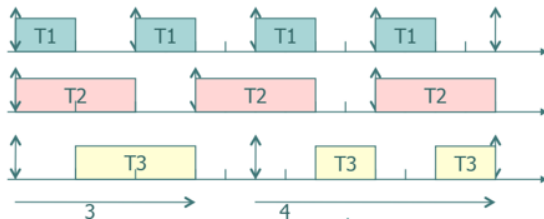
- $(P_1 = 4, C_1 = 2), (P_2 = 4, C_2 = 2), (P_3 = 12, C_3 = 7)$





# Scheduling anomalies

- Critical instant not necessarily the simultaneous release of higher priority tasks
  - $n=3, m=2$
  - $(P_1 = 2, C_1 = 1), (P_2 = 3, C_2 = 2), (P_3 = 4, C_3 = 2)$
  - Under RM scheduling
    - Response time of T3 higher at time 4 than at time 0

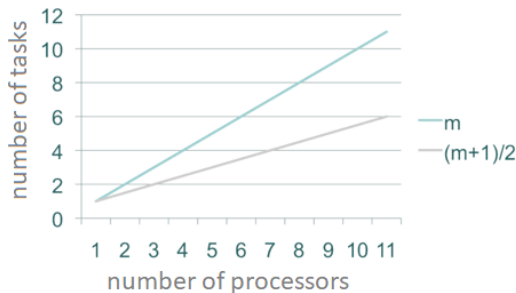


# General properties of multiprocessor scheduling

- Exact schedulability condition
  - $U \leq m$  and  $u_{max} \leq 1$ 
    - $U$  = Total utilization
    - $u_{max}$  = Maximum utilization
  - Does not tell for which scheduling algorithm!
- Schedule is cyclic on the hyperperiod  $H$  (PPCM( $P_i$ )) for:
  - Deterministic tasks
  - Without memory scheduling algorithms

# General properties of multiprocessor scheduling

- Theorem [Srinivasan Baruah, 2002]
  - Non existence of FPJ (FPJ+FPT) scheduling with utilization bound strictly larger than  $\frac{m+1}{2}$  for implicit deadline periodic task sets!



# Global multiprocessor scheduling: detailed outline

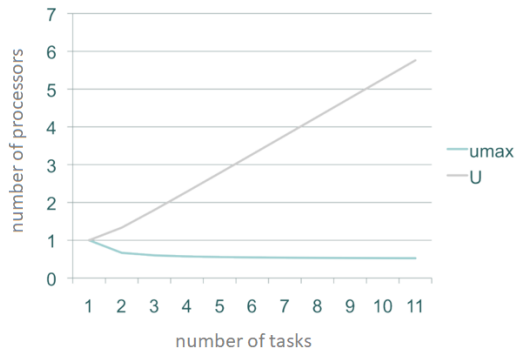
- Transposition of uni-processor algorithms
- Extensions of uni-processor algorithms
  - US (Utilization Threshold)
  - EDF(k)
  - ZL (Zero Laxity)
- Pfair approaches (Proportional Fair)

# Transposition of uni-processor algorithms

- Main algorithms
  - RM (Rate Monotonic)  $\rightarrow$  G-RM, Global RM
  - EDF (Earliest Deadline First)  $\rightarrow$  G-EDF, Global EDF
  - Not optimal anymore
  - Sufficient schedulability tests (depend on  $u_{max}$ )

G-RM	G-EDF
$u_{max} \leq m/(3m-2)$ and $U \leq m^2/(3m-2)$	$u_{max} \leq m/(2m+1)$ and $U \leq m^2/(2m+2)$
$u_{max} \leq 1/3$ and $U \leq m/3$	$u_{max} \leq 1/2$ and $U \leq (m+1)/2$
$U \leq m/2 * (1-u_{max}) + u_{max}$	$U \leq m - (m-1) u_{max}$

# Transposition of uni-processor algorithms



# Exten. of global RM/EDF: US (Utilization Threshold)

- Priority assignment depend on an utilization threshold  $\xi$ 
  - If  $u_{max} > \xi$  then  $T_i$  is assigned maximal priority
    - Else,  $T_i$ 's priority assigned as in original algorithm (RM/EDF)
- Remarks
  - Still non optimal
  - Outperforms the base policy that is used
  - Defies Dhall's effect

# Exten. of global RM/EDF: US (Utilization Threshold)

- Example: RM-US[ $\xi=\frac{1}{2}$ ]

	<b>C<sub>i</sub></b>	<b>P<sub>i</sub></b>	<b>U<sub>i</sub></b>	<b>Prio</b>
T1	4	10	2/5	2
T2	3	10	3/10	2
T3	8	12	2/3	$\infty$
T4	5	12	5/12	1
T5	7	12	7/12	$\infty$



# Exten. of global RM/EDF: US (Utilization Threshold)

- Utilization bounds

RM-US		EDF-US	
$\xi = m/(3m-2)$	$U \leq m^2/(3m-2)$	$\xi = m/(2m-1)$	$U \leq m^2/(2m-1)$
$\xi = 1/3$	$U \leq (m+1)/3$	$\xi = 1/2$	$U \leq (m+1)/2$

- Remarks

- Utilization bounds do not depend on  $u_{max}$  more
- EDF-US $[\xi = \frac{1}{2}]$  attains the best utilization bound possible for FPJ ( $\frac{m+1}{2}$ )

# Exten. of global RM/EDF: EDF(k)

- Task indices by decreasing utilization
  - $u_i \geq u_{i+1}$  for all  $i$  in  $[1, n]$
- Priority assignment depends on a threshold on task index
  - $i < k$ , then maximum priority
  - Else, priority assignment according to original algorithm

# Exten. of global RM/EDF: EDF(k)

- Example, EDF(4)

	<b>C<sub>i</sub></b>	<b>P<sub>i</sub></b>	<b>U<sub>i</sub></b>	<b>Prio</b>
T1	4	10	2/5	EDF
T2	3	10	3/10	EDF
T3	8	12	2/3	$\infty$
T4	5	12	5/12	$\infty$
T5	7	12	7/12	$\infty$

# Exten. of global RM/EDF: EDF(k)

- Sufficient schedulability test

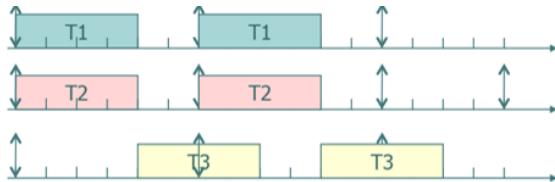
- $m \geq (k-1) - \lceil \frac{\sum_{i=k+1}^n u_i}{1-u_k} \rceil$
- $k_{min}$  = value minimizing right side of the equation
- With  $k = k_{min}$  , utilization bound of  $\frac{m+1}{2}$  (the best possible for FPJ)
- Comparison with EDF[ $\xi = \frac{1}{2}$ ]
  - Same utilization bound
  - EDF( $k_{min}$ ) dominates EDF[ $\xi = \frac{1}{2}$ ]

# Exten. of global RM/EDF: ZL (Zero Laxity) policies

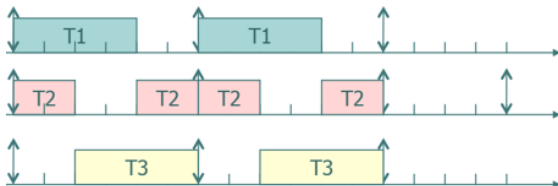
- XX-ZL: Apply policy XX until Zero Laxity
  - Maximal priority when laxity reaches zero (regardless of the currently running job), original priority assignment for the others
  - In category DPJ (dynamic job scheduling)
  - Policies: EDZL [Lee, 1994], RMZL [Kato et al, 2009], FPZL [Davis et al, 2010]
  - Utilization bound:  $\frac{m+1}{2}$
  - Dominates G-EDF

# Exten. of global RM/EDF: ZL (Zero Laxity) policies

- Example:  $n=3, m=2$ ; all  $P_i$  to 6, all  $C_i$  to 4
- G-EDF: T3 misses its deadline



- EDZL: OK



# Pfair algorithms: principle (up to/from here: session 24/25)

- Pfair: “Proportionate Fair” [Baruah et al, 1996]
  - Allocate time slots to tasks as close as possible to a “fluid” system, proportional to their utilization factor
- Example
  - $C_1=C_2=3, P_1=P_2=6$  ( $u_1=u_2=\frac{1}{2}$ )
  - Each task will be “approximately” allocated 1 slot out of 2 (whatever the processor)

# Pfair algorithms: principle

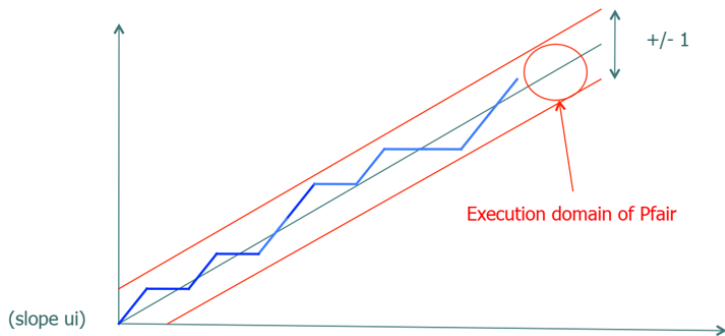
- Lag function: Difference between real and fluid execution
  - Discrete time, successive time slots  $[t, t+1]$
  - Weight of a task:  $\omega_i = u_i$
- Lag
  - $lag(T_i, t) = \omega_i t - \sum_{u=0}^{t-1} S(T_i, u)$
  - First term: Fluid execution
  - Second term: real execution, with  $S(T_i, u) = 1$  if  $T_i$  executed in slot  $u$ , else 0

**Pfair schedule: for all time  $t$ , lag in interval  $[-1, 1]$**



# Pfair algorithms: principle

- Example



# Pfair algorithms: principle

- Property
  - If a Pfair schedule exists, deadlines are met
- Exact test of existence of a Pfair schedule
  - $\sum_{i=1}^n u_i < m$

**Full processor utilization!**

# Pfair algorithms: construction of a Pfair schedule

- Divide tasks in unity-length sub-tasks
  - Pfair condition: each subtask  $j$  executes in a time window between a pseudo-arrival and a pseudo deadline
  - Pseudo-arrival
    - $r(T_i^j) = \lfloor \frac{j-1}{\omega_i} \rfloor$
  - Pseudo-deadline
    - $d(T_i^j) = \lceil \frac{j}{\omega_i} \rceil$

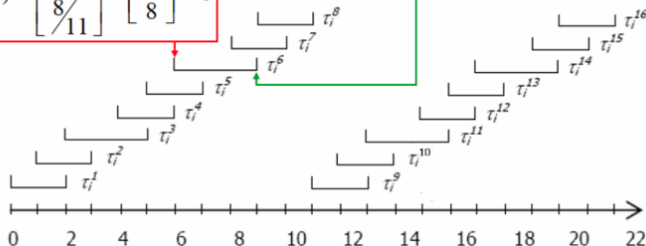
# Pfair algorithms: construction of a Pfair schedule

## • Example (to be fixed)

■  $(C_i = 8, T_i = 11) \rightarrow \omega_i = u_i = 8/11$

$$r(\tau_i^6) = \left\lfloor \frac{6-1}{8/11} \right\rfloor = \left\lfloor \frac{55}{8} \right\rfloor = 6$$

$$d(\tau_i^6) = \left\lceil \frac{6}{8/11} \right\rceil = \left\lceil \frac{33}{4} \right\rceil = 9$$



# Pfair algorithms: scheduling algorithms

- EPDF (Earliest Pseudo-Deadline First)
  - Apply EDF to pseudo-deadlines
  - Optimal only for  $m=2$  (2 processors)
- Ongoing works
  - Reduce numbers of context switches and migrations while maintaining optimality

# Conclusion

- Multi-processor scheduling is an active research area
- Ongoing works
  - Global multi-core scheduling
  - Semi-partitioned scheduling
  - Determining upper bounds of practical factors (preemption, migration, ...)
  - Implementation in real-time operating systems

*The End*